

Evaluating Classifiers

Machine Learning

PHYS 453 – Spring 2022

Dr. Daugherty

Evaluating Classifiers

- How can I measure how well a classifier works?
- Where do I look for ways to improve performance?

Sources:

- https://scikit-learn.org/stable/modules/model_evaluation.html#classification-metrics
- Binary Classification Metrics paper, on canvas or: <https://arxiv.org/pdf/1410.5330>

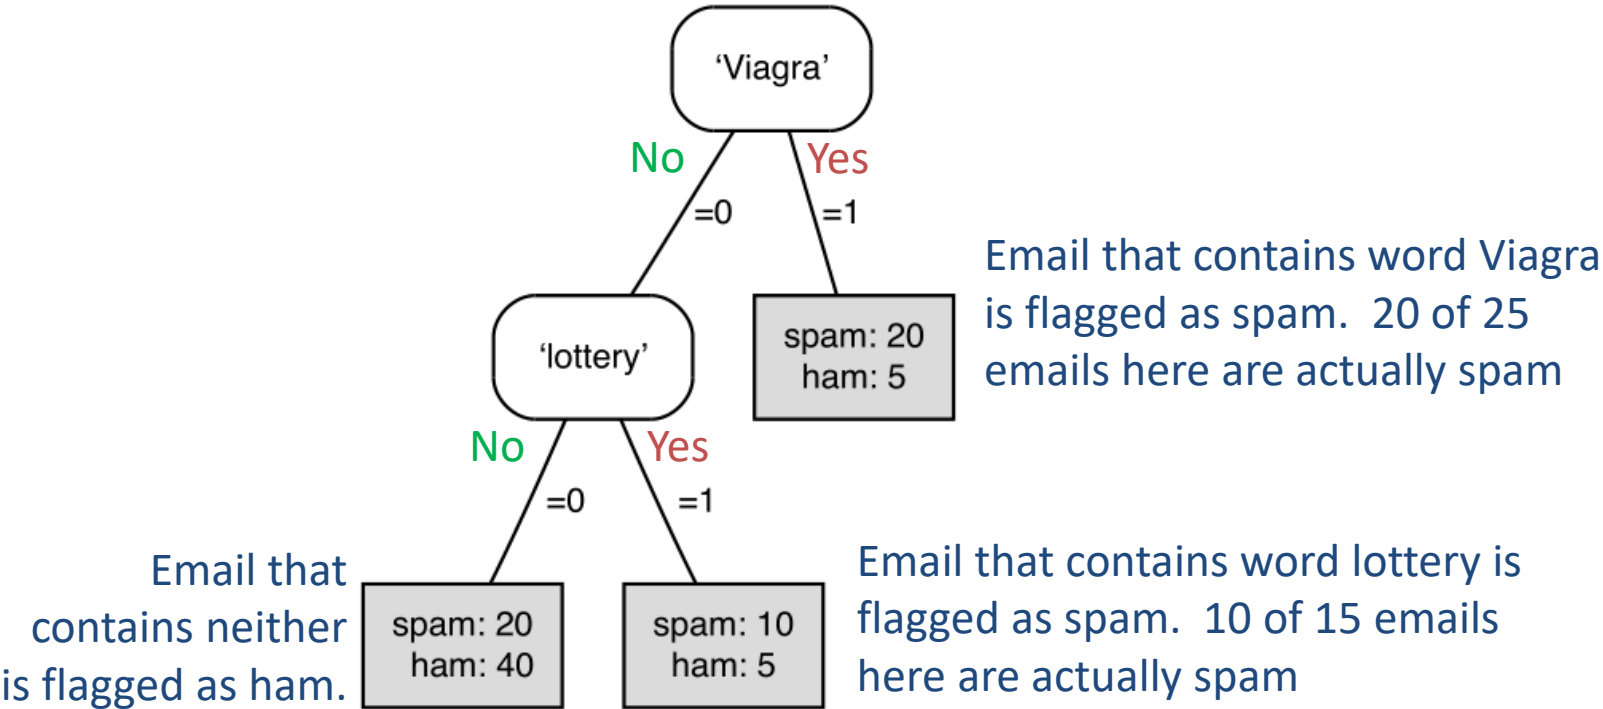
Evaluating Classifiers

CONFUSION MATRIX

Confusion Matrix

		Predicted class	
		<i>P</i>	<i>N</i>
Actual Class	<i>P</i>	True Positives (TP)	False Negatives (FN)
	<i>N</i>	False Positives (FP)	True Negatives (TN)

Spam detection decision tree



	SPAM Predicted ⊕	HAM Predicted ⊖	
Actual ⊕	30	20	50
Actual ⊖	10	40	50
	40	60	100

Accuracy Metrics

		Predicted class	
		<i>P</i>	<i>N</i>
Actual Class	<i>P</i>	True Positives (TP)	False Negatives (FN)
	<i>N</i>	False Positives (FP)	True Negatives (TN)

$$ERR = \frac{FP + FN}{FP + FN + TP + TN} = 1 - ACC$$

Error %

$$ACC = \frac{TP + TN}{FP + FN + TP + TN} = 1 - ERR$$

Accuracy %

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN}$$

False Positive Rate = (# of FP) / (# actually N)
“what percentage of the real N did I miss?”

$$TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

True Positive Rate = (# of TP) / (# actually P)
“what percentage of the real P did I get?”

$$PRE = \frac{TP}{TP + FP}$$

PRECISION = the ability of the classifier not to label as positive a sample that is negative.
Fraction of pos guesses that are right.

$$REC = TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

RECALL = the ability of the classifier to find all the positive samples

$$F_1 = 2 \cdot \frac{PRE \cdot REC}{PRE + REC}$$

Fraction of all actual pos we guessed as pos.

F1 Score = combines both into a single number. 1 is perfect.

Challenge: gotta find them all!

	Predicted \oplus	Predicted \ominus	
Actual \oplus	30	20	50
Actual \ominus	10	40	50
	40	60	100

		Predicted class	
		P	N
Actual Class	P	True Positives (TP)	False Negatives (FN)
	N	False Positives (FP)	True Negatives (TN)

$$ERR = \frac{FP + FN}{FP + FN + TP + TN} = 1 - ACC$$

$$ACC = \frac{TP + TN}{FP + FN + TP + TN} = 1 - ERR$$

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN}$$

$$TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

$$PRE = \frac{TP}{TP + FP}$$

$$REC = TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

$$F_1 = 2 \cdot \frac{PRE \cdot REC}{PRE + REC}$$

Measure	Definition	Equal to	Estimates
number of positives	$Pos = \sum_{x \in Te} I[c(x) = \oplus]$		
number of negatives	$Neg = \sum_{x \in Te} I[c(x) = \ominus]$	$ Te - Pos$	
number of true positives	$TP = \sum_{x \in Te} I[\hat{c}(x) = c(x) = \oplus]$		
number of true negatives	$TN = \sum_{x \in Te} I[\hat{c}(x) = c(x) = \ominus]$		
number of false positives	$FP = \sum_{x \in Te} I[\hat{c}(x) = \oplus, c(x) = \ominus]$	$Neg - TN$	
number of false negatives	$FN = \sum_{x \in Te} I[\hat{c}(x) = \ominus, c(x) = \oplus]$	$Pos - TP$	
proportion of positives	$pos = \frac{1}{ Te } \sum_{x \in Te} I[c(x) = \oplus]$	$Pos/ Te $	$P(c(x) = \oplus)$
proportion of negatives	$neg = \frac{1}{ Te } \sum_{x \in Te} I[c(x) = \ominus]$	$1 - pos$	$P(c(x) = \ominus)$
class ratio	$clr = pos/neg$	Pos/Neg	
(*) accuracy	$acc = \frac{1}{ Te } \sum_{x \in Te} I[\hat{c}(x) = c(x)]$		$P(\hat{c}(x) = c(x))$
(*) error rate	$err = \frac{1}{ Te } \sum_{x \in Te} I[\hat{c}(x) \neq c(x)]$	$1 - acc$	$P(\hat{c}(x) \neq c(x))$

A slightly different version of the same thing, just in case...

Measure	Definition	Equal to	Estimates
true positive rate, sensitivity, recall	$tpr = \frac{\sum_{x \in Te} I[\hat{c}(x)=c(x)=\oplus]}{\sum_{x \in Te} I[c(x)=\oplus]}$	TP/Pos	$P(\hat{c}(x) = \oplus c(x) = \oplus)$
true negative rate, specificity	$tnr = \frac{\sum_{x \in Te} I[\hat{c}(x)=c(x)=\ominus]}{\sum_{x \in Te} I[c(x)=\ominus]}$	TN/Neg	$P(\hat{c}(x) = \ominus c(x) = \ominus)$
false positive rate, false alarm rate	$fpr = \frac{\sum_{x \in Te} I[\hat{c}(x)=\oplus, c(x)=\ominus]}{\sum_{x \in Te} I[c(x)=\ominus]}$	$FP/Neg = 1 - tnr$	$P(\hat{c}(x) = \oplus c(x) = \ominus)$
false negative rate	$fnr = \frac{\sum_{x \in Te} I[\hat{c}(x)=\ominus, c(x)=\oplus]}{\sum_{x \in Te} I[c(x)=\oplus]}$	$FN/Pos = 1 - tpr$	$P(\hat{c}(x) = \ominus c(x) = \oplus)$
precision, confidence	$prec = \frac{\sum_{x \in Te} I[\hat{c}(x)=c(x)=\oplus]}{\sum_{x \in Te} I[\hat{c}(x)=\oplus]}$	$TP/(TP + FP)$	$P(c(x) = \oplus \hat{c}(x) = \oplus)$

Table : A summary of different quantities and evaluation measures for classifiers on a test set Te . Symbols starting with a capital letter denote absolute frequencies (counts), while lower-case symbols denote relative frequencies or ratios. All except those indicated with (*) are defined only for binary classification.

Suppose a classifier's predictions on a test set are as in the following table:

	Predicted \oplus	Predicted \ominus	
Actual \oplus	60	15	75
Actual \ominus	10	15	25
	70	30	100

From this table, we see that the true positive rate is $tpr = 60/75 = 0.80$ and the true negative rate is $tnr = 15/25 = 0.60$. The overall accuracy is $acc = (60 + 15)/100 = 0.75$, which is no longer the average of true positive and negative rates. However, taking into account the proportion of positives $pos = 0.75$ and the proportion of negatives $neg = 1 - pos = 0.25$, we see that

$$acc = pos \cdot tpr + neg \cdot tnr$$

<https://xkcd.com/2236/>

IS IT CHRISTMAS?

|<

< PREV

RANDOM

NEXT >

>|

NO*

*99.73% ACCURATE

XKCD.COM PRESENTS A NEW "IS IT CHRISTMAS"
SERVICE TO COMPETE WITH ISITCHRISTMAS.COM

|<

< PREV

RANDOM

NEXT >

>|

sklearn.metrics.confusion_matrix

`sklearn.metrics.confusion_matrix(y_true, y_pred, *, labels=None, sample_weight=None, normalize=None)`

[\[source\]](#)

Compute confusion matrix to evaluate the accuracy of a classification.

By definition a confusion matrix C is such that $C_{i,j}$ is equal to the number of observations known to be in group i and predicted to be in group j .

Thus in binary classification, the count of true negatives is $C_{0,0}$, false negatives is $C_{1,0}$, true positives is $C_{1,1}$ and false positives is $C_{0,1}$.

Read more in the [User Guide](#).

Parameters: **y_true** : *array-like of shape (n_samples,)*

Ground truth (correct) target values.

y_pred : *array-like of shape (n_samples,)*

Estimated targets as returned by a classifier.

labels : *array-like of shape (n_classes), default=None*

List of labels to index the matrix. This may be used to reorder or select a subset of labels. If `None` is given, those that appear at least once in `y_true` or `y_pred` are used in sorted order.

sample_weight : *array-like of shape (n_samples,), default=None*

Sample weights.

New in version 0.18.

normalize : *{'true', 'pred', 'all'}, default=None*

Normalizes confusion matrix over the true (rows), predicted (columns) conditions or all the population. If `None`, confusion matrix will not be normalized.

```
>>> from sklearn.metrics import confusion_matrix
>>> y_true = [2, 0, 2, 2, 0, 1]
>>> y_pred = [0, 0, 2, 2, 0, 2]
>>> confusion_matrix(y_true, y_pred)
array([[2, 0, 0],
       [0, 0, 1],
       [1, 0, 2]])
```

```
>>> y_true = ["cat", "ant", "cat", "cat", "ant", "bird"]
>>> y_pred = ["ant", "ant", "cat", "cat", "ant", "cat"]
>>> confusion_matrix(y_true, y_pred, labels=["ant", "bird", "cat"])
array([[2, 0, 0],
       [0, 0, 1],
       [1, 0, 2]])
```

```
>>> from sklearn.metrics import classification_report
>>> y_true = [0, 1, 2, 2, 2]
>>> y_pred = [0, 0, 2, 2, 1]
>>> target_names = ['class 0', 'class 1', 'class 2']
>>> print(classification_report(y_true, y_pred, target_names=target_names))
```

	precision	recall	f1-score	support
class 0	0.50	1.00	0.67	1
class 1	0.00	0.00	0.00	1
class 2	1.00	0.67	0.80	3
accuracy			0.60	5
macro avg	0.50	0.56	0.49	5
weighted avg	0.70	0.60	0.61	5

Overall accuracy = 0.60



Chapter 2.1

Difficulty: 2

There are 20 dogs(+) and 10 cats(-). A binary classifier correctly predicts 5 dogs and incorrectly predicts 5 cats. Fill in the following contingency for this binary classifier matrix.

	Predicted +	Predicted -	
Actual +			
Actual -			

Chapter 2.1

Difficulty: 2

Given that:

Total = 100

False Negatives = 10

Precision = $\frac{4}{5}$

Recall = $\frac{6}{7}$

Can you complete the contingency matrix.

	Predicted +	Predicted -	
Actual +			
Actual -			

Chapter 2.1

Difficulty: 4

Two binary classifiers are used to predicted whether a patent has a life threatening diseases or not. Decide whether Classifier A or B would be better at reducing casualties.

A

10	10	20
40	9940	9980
50	9950	10000

B

13	7	20
87	9813	9980
100	9900	10000

Two contingency matrices.

Tutorials

- <https://github.com/mdaugherty/PatternRecognition2018/blob/master/Tutorial%203-1.ipynb>
- <https://github.com/mdaugherty/PatternRecognition2018/blob/master/Tutorial%203-2.ipynb>

Summary

Know the following:

- Accuracy / error rate
- TP, FP, TN, FN in confusion matrix
- Precision
- Recall